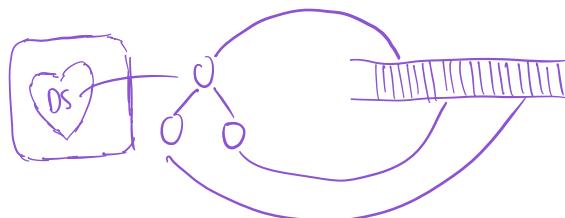


# MEMORY MANAGEMENT

## WHY MEM MGMT in DS?

- see "Motivation to DS" again



## WHAT exactly is MEM MGMT?

- overriding C++ new/delete

0max ...

lagg ...

crash ...

?!!! ing C++ new/delete ???

THE MS QNS:

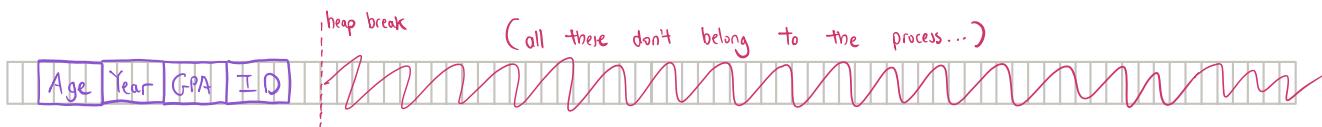
WHY do we want to reinvent the

THE M ANS:  
consider the struct Student {

Q: WHAT is sizeof(Student)? (assume 32-bit WIN)  
A: 4 bytes \* 4 = 16 bytes

```
int Age;  
long Year;  
float GPA;  
long ID;
```

```
Student* s1 = new Student();
```



```
Student* s2 = new Student();
```

this will trigger...

① allocate() from proc heap

- find empty slot
- reduce fragmentation
- etc...

② OUT\_OF\_MEM\_ERR

③ SYS\_CALL\_TRAP

only SYS has oversight across procs

(Context switch) USER MODE

KERNEL MODE

④ lock()

- prevent 2 procs alloc same mem

⑤ allocate() more mem to proc



⑥ BUT maybe unlucky... hit OUT\_OF\_RAM\_ERROR

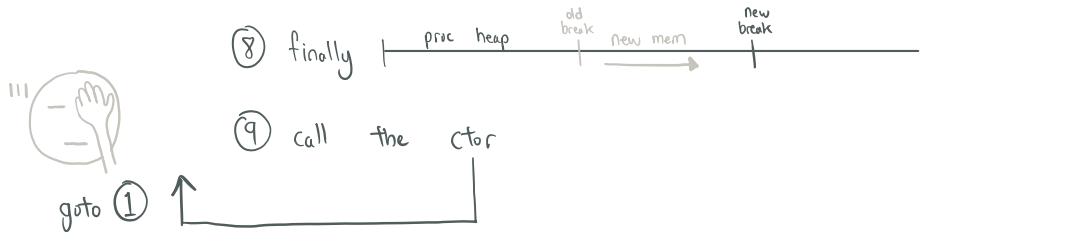


⑦ swap old/unused pages to HDD



tl;dr

THIS is very BAD



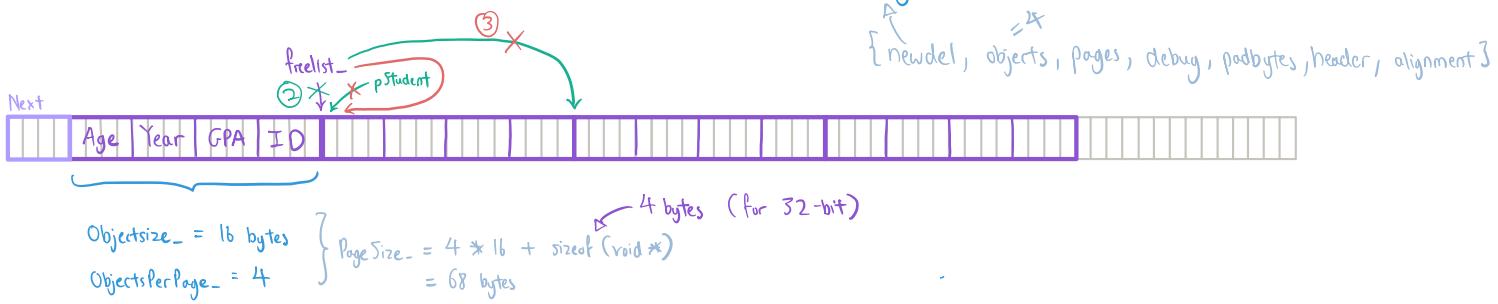
this is WHY we have ...

## CUSTOM ALLOCATORS

MAIN IDEA:

- ① alloc a "chunk" (a.k.a. pages) of mem at one go

ObjectAllocator\* oa = new ObjectAllocator(sizeof(Student), config);



- ② when user `Student* pStudent = static_cast<Student*>(oa->Allocate());`

return 1 block from a page

instead of calling `new Student();`

- ③ when user `oa->Free(pStudent)`

"put" back the block into the free list

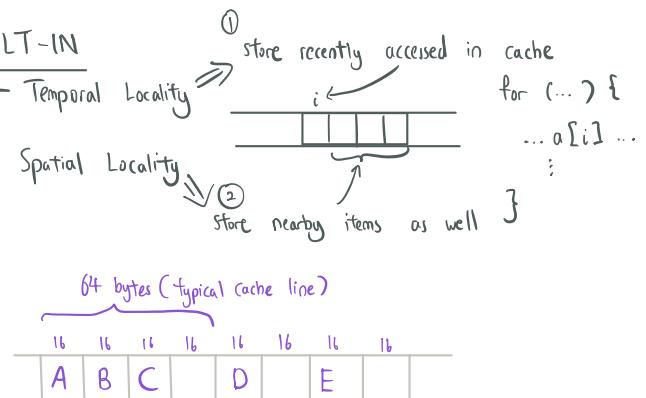
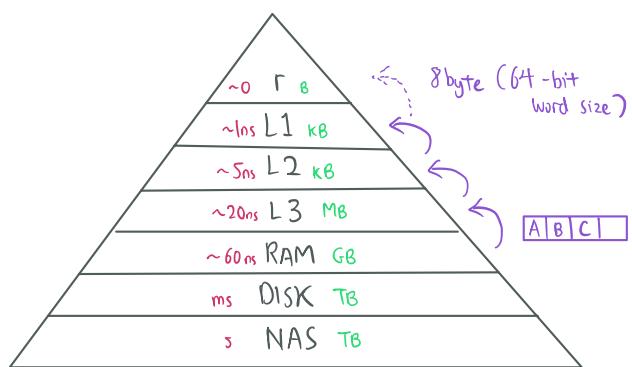
- ④ ... reuse blocks by adhering to ② and ③

Q: Show what happens to the ptrs in mem.

now let's look deeper at

## WHY CUSTOM OAS CAN BE BETTER THAN BUILT-IN

- ## ① performance boost with LOCALITY OF REFERENCE

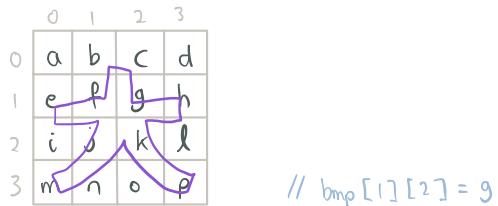


Q: Describe what happens when

do(A), do(C), do(A), do(B), do(D)



now consider this bmp operation ...



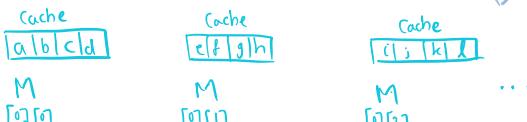
Pixel bmp[4][4]; // sizeof(Pixel) = 16, cache line = 64

```
for (auto i=0; i<4; ++i)
```

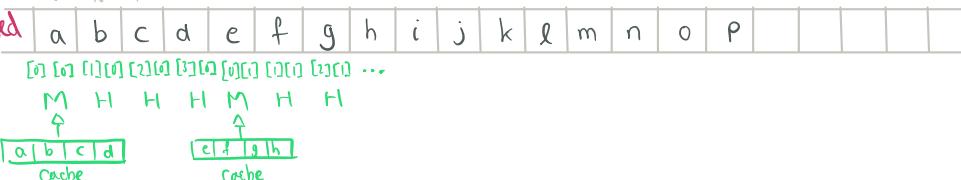
```
for ( auto j=0; j<4; ++j )
```

`clear(bmp[i][j]); //①` vs `clear(bmp[j][i]); //②`

Q: Does it matter if we choose  or ?



xx



and time for Misses >> Hits

(much  
slower)

$\Rightarrow$  ~~graph~~ is better

- ② customized to specific app needs



### KNOWN OBJ SIZES

- fixed entity types
- max objs in scene

## ALLOC PATTERNS

- Same sequence between levels
- fixed durations

### ③ MEM DEBUGGING

- count allocs /deallocs (detect leaks)
  - track alloc sizes
  - time between allocs
  - etc...

So what are some common

## CUSTOM ALLOCATOR TYPES ?

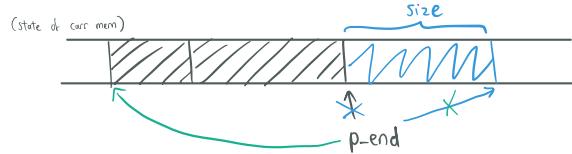


### LINEAR ALLOCATOR

oa → Allocate (size);

oa → FreeAll();

Q: Show what happens in mem after the above stnts are ran...



✓ fast

✓ insignificant space overhead

✗ no individual frees

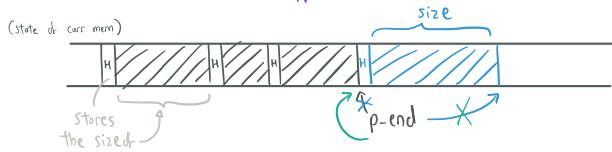
} Suitable for e.g. glx screen buffer  
- clrScr() every frame

### STACK ALLOCATOR

oa → Allocate (size);

oa → Free ();

Q: Show what happens in mem after the above stnts are ran...



✓ still fast

✗ slight space overhead

✓ can individually Free() BUT only LIFO

} Suitable for e.g. hierarchical game levels - 02 - 02.1 - 02.2

### POOL ALLOCATOR

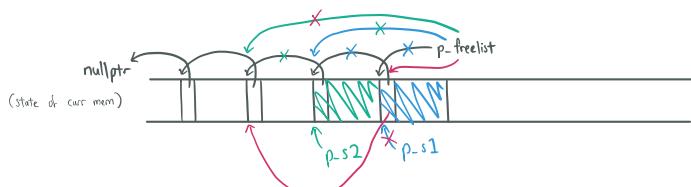
after alloc "chunk", pre-split into blocks, ie. populate the free list

Student\* p\_s1 = static\_cast<Student\*>(oa → Allocate());

Student\* p\_s2 = static\_cast<Student\*>(oa → Allocate());

oa → Free(p\_s1); // return to free list for reuse

Q: Show what happens in mem after the above stnts are ran...



✓ still fast

✓ fixed-size blocks

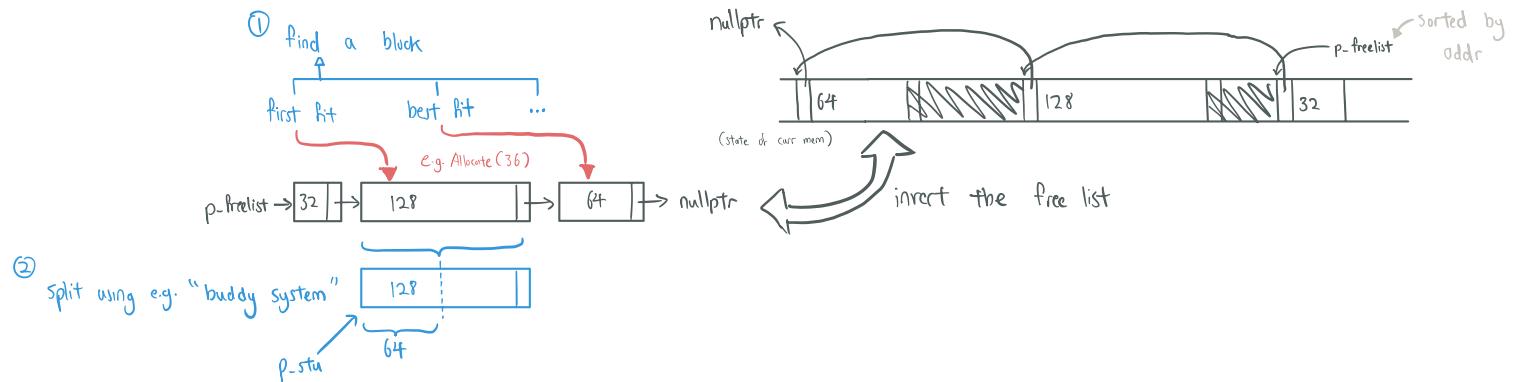
✗ some space overhead

} Suitable for e.g. game entities  
- roughly same size  
- frequent arbitrary alloc/frees

## FREE LIST ALLOCATOR (a.k.a. pooling with dynamic sizing)

(ignoring the cuts...)

$p\_stu = \text{oa} \rightarrow \text{Allocate}(\text{size})$ ;

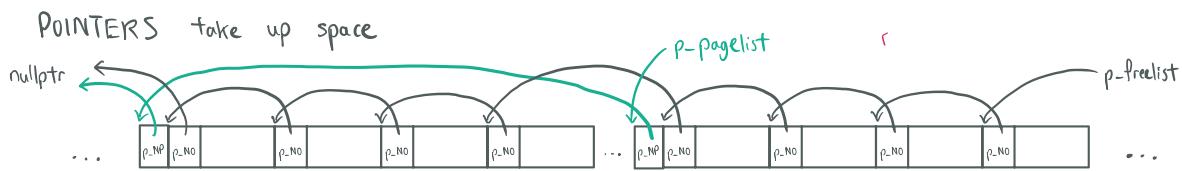


oa → Free( $p\_stu$ ) ① link back to free list (in order of address)  
② merge if contiguous (coalescence)

✓ general purpose	Suitable for ANYTHING - BUT always consider more specialized list - we can also note that as it gets more generic, speed decreases.
✗ not as fast	
(still faster than malloc)	

visit [github.com / mtrebi / memory-allocators](https://github.com/mtrebi/memory-allocators) to see some benchmarks of the above

## MEM ALLOCATOR OVERHEAD



STATISTICS take up space

OAStats →	Freeobjects_-	ObjectsInuse_-	PageInuse_-	MostObjects_-	Allocations_-	Deallocations_-
◇						

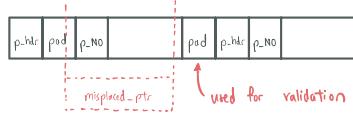
### DEBUGGING

- e.g. if ( $\text{Allocations}_- == \text{Deallocations}_-$ ) noLeaks = true;

### IMPROVE OA DESIGN

- e.g. if ( $\text{MostObjects}_- > \text{threshold}$ )  $\text{++ ObjectsPerPage}_-$ ;

and then padding (for validation)



and then there're HEADERS...



ALIGNMENT also takes up space

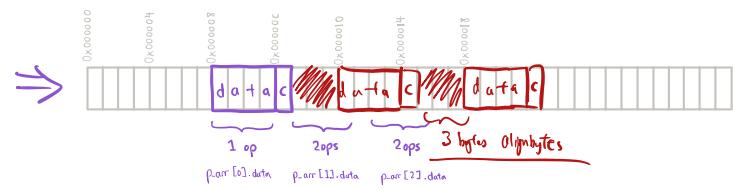
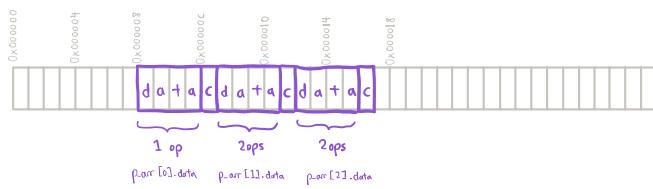
e.g. struct Intcy {  
    int data; //4  
    char c; //1  
};

Intcy\* p-arr = new Intcy[3];

Q: Show the above array in mem below... (assume 32-bit archi)

Q: How many atomic operations to access each element of the array?

$$\rightarrow \frac{32}{8 \text{ bit}} = 4 \text{ byte word}$$



Q: What is the sol<sup>n</sup> to this problem?

A: Add alignbytes until data within boundaries.

# lastly, a note about FRAGMENTATION

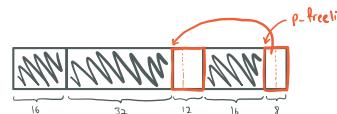
2 types

e.g. a fixed-size pool allocator  
of 16-byte blocks

e.g. a fixed-size pool allocator  
of 16-byte blocks

## EXTERNAL

e.g. a freelist allocator



What if now we want

`p_0 = static_cast<sixteenByte*>(oa->Allocate());`

(fragment is in free men)

(fragment is within allocated mem)

